



A positive signal, measured twice

The required unique-user KPI shows a 13.9% lift. A user-cohort sensitivity view is smaller, near 5.5%, because assignment can change across months.



Decision statement

Treatment conversion was **2.23%** versus **1.95%** for control. The conventional two-proportion test gives **p = 3.42e-09**.



The simple control-rate counterfactual estimates **1,569 additional conversions**.

AUDIT NOTE

26,426 users changed groups across months. The user-cohort sensitivity estimate is +5.5% with p = 0.016.

Evidence by cohort, device, and analysis unit

Treatment stayed above control in every cohort, while the effect size depends on how repeated users are represented.



Robustness table


Analysis unit	Treatment CR	Control CR	Relative uplift	p-value
Brief-defined unique users	2.23%	1.95%	+13.9%	3.42e-09
User-cohort	1.99%	1.89%	+5.5%	0.0162
Session	1.98%	1.88%	+5.6%	0.0142

What this means

- All five monthly cohorts show positive brief-defined uplift, ranging from 11.6% to 27.6%.
- Both experiment arms declined from February to June, suggesting a material time effect.
- Desktop shows the strongest descriptive device uplift at 29.1%.

Method, data controls, and next decision

The portfolio version prioritizes reproducibility, privacy, and honest experiment interpretation.



Method

- Remove invalid cohort 203902.
- Remove exact duplicate source rows.
- Fill missing device and gender categories.
- Calculate brief-defined user CR and VPT.
- Run a two-sided proportions test and confidence intervals.
- Compare user, user-cohort, and session analysis units.

Privacy controls

The public project excludes raw CSV files, userhash, sessionid, proprietary assignment screenshots, passwords, and internal handoff documents. Published data tables are aggregate only.



RECOMMENDATION

Treat the experiment as a positive signal. Confirm the randomization unit before rollout sizing, investigate the cohort decline, and use the Desktop result as a follow-up hypothesis rather than a final targeting rule.